

VoltChain Technical Reference

A mission-critical AI intelligence dashboard providing sub-millisecond predictive maintenance and unified supply chain visibility.

99.7%

UPTIME TARGET

<2ms

CORE LATENCY

67

TESTS PASSING

12.4t+

CO2 OFFSETS

Author: Saurabh Kumar

Repository: github.com/SaurabhForge/VoltChain

Live Demo: saurabhforge-voltchain-frontend.onrender.com

[SUBMISSION REFERENCE](#)

SYSTEM OVERVIEW & DATA PIPELINE

VoltChain uses a decoupled, two-tier architecture designed to run on lightweight, cost-effective infrastructure. The front-end consists of a highly optimized static SPA served via CDN, while the back-end is powered by Express 5 running on Node.js 20. Telemetry data is streamed continuously using **Server-Sent Events (SSE)**, providing an open, lightweight, unidirectional socket connection with native auto-reconnect behavior.

```
[CDN Static Frontend] <—(SSE Telemetry /api/telemetry/stream)— [Express 5 API Service]
[SPA Client Controllers] —(REST Fetch: Stats/Timeline/QMS)—> [In-Memory State Engine]
[D3.js Graph & Three.js] <—(REST Fetch: Graph Nodes)—> [Isolated State Cloner]
```

Infrastructure as Code: Render Blueprint configuration (`render.yaml`) manages the deployment. It enforces CORS constraints (restricting the API strictly to the frontend origin) and injects essential security headers (such as `X-Frame-Options: DENY` and `X-Content-Type-Options: nosniff`).

THE "BRAINS" BEHIND VOLTCHAIN

MCP AI Analyst AI INFERENCE

The AI analyst processes real-time telemetry to forecast SOH degradation curves, allowing for proactive, rather than reactive, maintenance scheduling. It utilizes template-driven natural language routing to supply instant context-aware operational advice, routing 15 key categories (such as battery temperature, alerts, charging schedules, and compliance risks) with sub-15ms response times.

Neural Knowledge Graph D3 ENGINE

The Knowledge Graph maps 2.4B simulated supply chain entities to identify supply chain bottlenecks before they ripple through the fleet. Rendered dynamically on the frontend via a **D3 v7 force-directed simulation**, it visualizes real-time cargo relationships and risk scores (Low, Medium, High risk indicators) directly from mine to gigafactory.




Fig 2.1 – Live Command Dashboard showing active SVG telemetry charts and the embedded MCP AI Analyst

COMPREHENSIVE TESTING SUITE

VoltChain is backed by **67 automated tests** executing sequentially under the Vitest test runner (utilizing a single-fork model to prevent database state collision).

- **37 Unit Tests (`tests/unit.test.js`):** Validates chat keyword routing, statistical bounds of random-walk mutations, QMS defect flags, input constraints, and per-connection state isolation.
- **30 Integration Tests (`tests/api.integration.test.js`):** Uses Supertest to verify all 9 API routes, CORS preflight policies, 404 behavior, and parameter validation.


STRESS & LOAD TESTING

To verify the platform's stability, the backend was subjected to high-throughput load tests using ****Autocannon****.

- **REST Endpoint Saturation:** Under a load of 100 concurrent connections pipelined by 10, the telemetry API achieved a throughput of **3,632 requests per second** with **0 errors**.
- **POST Chat Saturation:** Under 20 concurrent connections, the chat endpoint maintained an average throughput of **1,248 requests per second** with a p50 latency of only **15ms** and **0 errors**.
- **SSE Client Scaling:** The SSE stream was tested with 50 concurrent persistent clients for 20 seconds, pushing telemetry updates cleanly without packet loss or socket leaks.

RESILIENCE & FAILOVER MECHANISMS

Resilience Layer	Failure Condition	System Behavior
Graceful Degradation	API service goes offline	Frontend switches to client-side backup simulation. Chat widget continues working locally, D3 Graph falls back to a 4-node model, and metrics continue to update.
Crash Guards	Uncaught exceptions / unhandled promise rejections	Process-level event listeners trap the error, write a structured JSON error log, and perform a clean exit, preventing silent thread hangs.
Health Monitoring	Render platform health checks	Dedicated <code>/health</code> endpoint exposes live memory footprint, uptime, and timestamp in under 1ms. Structured JSON logs stream to stdout.
State Isolation	Concurrent SSE connections	Each SSE client has their state isolated via per-connection cloning, resolving cross-client variable pollution.



ALIGNMENT: ACCELERATING NET ZERO

The production of a single EV battery packs a significant environmental footprint, generating substantial greenhouse gas emissions before the vehicle even drives its first mile. Extending the operational lifespan of these batteries is a key lever in reducing global industrial carbon footprints.

VoltChain's MCP AI Analyst directly tackles this challenge by moving fleet dispatchers from reactive maintenance to proactive preservation. By tracking variables such as cathode temperature, internal resistance, and cycle counts, the platform calculates continuous SOH degradation curves.

JUDGE'S EVALUATION FLOW (STEP-BY-STEP)

- 1. Access the Dashboard:** Scroll down to the *Live Command Dashboard* section at saurabhforge-voltchain-frontend.onrender.com.
 - Observe the **Uptime status badge**: it will show **LIVE** if connected to the API, and **OFFLINE** if running client-side simulation.
 - Watch the KPI cards and the SVG chart. They will update every 2 seconds via the live telemetry feed.
- 2. Interact with the MCP AI Analyst:** Locate the Chat Widget in the dashboard overview.
 - Type a query such as `"Show current SOH"` or `"Any anomalies?"` and press Enter.
 - The analyst will display a typing indicator and return a response based on current metrics. The log output in your terminal will record the request as a structured JSON entry.
- 3. Scrub the Time Capsule:** Scroll down to the *Temporal Time Capsule*. Drag the slider from **TODAY** back to **Q1-2024**. Observe how the events log shifts to historical incidents, and see the dashboard's active vehicle counts change accordingly.
- 4. Examine the Neural Knowledge Graph:** Scroll to the bottom. Drag nodes to reshape the graph. Hover over any node to view its health score and compliance classification.




Fig 4.1 – Time Capsule showing historical Q1 snapshot with live event log

Note on Free Tier Services: Render's free tier spins down the backend API web service after 15 minutes of inactivity. When loading the site for the first time, the API may take 30-50 seconds to warm up. The frontend handles this delay gracefully, showing an "Offline" status before transitioning to "Live" once the backend responds.